

# Aravinda Raman Jatavallabha

Raleigh, NC | [aravindaraman14@gmail.com](mailto:aravindaraman14@gmail.com) | (919) 327-0958 | [aravinda-1402.github.io](https://aravinda-1402.github.io) | [linkedin.com/in/aravinda-jatavallabha](https://linkedin.com/in/aravinda-jatavallabha)

## EDUCATION

|  |                        |
|--|------------------------|
| <b>Master of Computer Science (Data Science Track)</b>   North Carolina State University, Raleigh, NC              | Aug 2023-May 2025      |
| <u>Courses</u> – Data Science, Natural Language Processing, Neural Networks, Database Management Systems           | <b>CGPA: 4.0/4.0</b>   |
| <b>B. Tech in Information Technology</b>   Manipal Institute of Technology, Manipal, India                         | Jun 2019-Jul 2023      |
| <u>Minor</u> : Big Data Analytics; <u>Courses</u> - Data Mining, Machine Learning, Pattern Recognition, Algorithms | <b>CGPA: 8.64/10.0</b> |

## TECHNICAL SKILLS

- **Programming Languages & Frameworks** : Python, SQL, TypeScript, JavaScript, Spring Boot, Angular, React, Flask, FastAPI, REST APIs
- **Tools & Platforms**: Docker, Git, Linux, Power BI, Azure OpenAI, AWS (S3, SageMaker, Lambda), Snowflake
- **Libraries**: Pandas, NumPy, Matplotlib, Scikit-learn, TensorFlow, Keras, PyTorch, LangChain, SpaCy, NLTK, SciPy, PyG
- **Machine Learning**: Time Series Analysis, Classification, Regression, Convolutional Neural Networks (CNN), Natural Language Processing (NLP), Graph Neural Networks (GNN), Retrieval-Augmented Generation (RAG), Large Language Models (LLMs), Prompt Engineering
- **Training & Certifications**: [Deep Learning \(deeplearning.ai\)](#), [Machine Learning \(Stanford Online\)](#), [AI Summer School](#)

## WORK EXPERIENCE

|  |                   |
|--|-------------------|
| <b>Machine Learning Engineer – Full Stack Co-op</b>   SmartProtect Public Safety Solutions, Wilmington, DE   | May 2024-Current  |
| <ul style="list-style-type: none"><li>• Developed and <b>A/B tested</b> time series <b>predictive models</b> (ARIMA, FB Prophet, LSTM) on <b>1.2M+ call records</b>; deployed real-time <b>FastAPI</b> inference endpoints that improved scheduling accuracy by 20% and reduced dispatcher wait time by 14%.</li><li>• Productionized <b>ML pipelines</b> using <b>Flask APIs</b>, <b>AWS SageMaker</b>, and <b>Snowflake</b>, cutting model retraining time by 35% via <b>CI/CD</b> orchestration; integrated <b>Azure OpenAI LLMs</b> for anomaly summarization and transcript Q&amp;A.</li><li>• Built <b>internal ops dashboard</b> using <b>Spring Boot + Angular</b>, powering live analytics for shift forecasting, <b>LLM-driven alerts</b>, and scheduling KPIs — used daily by 6+ teams across 3 regional call centers.</li><li>• Designed <b>optimization algorithms</b> for staff coverage using call volume clustering and anomaly tags, lowering overtime by 18% and boosting resource utilization by 22%.</li><li>• Implemented <b>model monitoring</b> and <b>data drift detection</b> using statistical checks, version tracking, and pipeline alerts — increasing post-deployment reliability by 40% and enabling auditability for compliance.</li></ul> |                   |
| <b>Machine Learning Engineer Intern</b>   Defence Research and Development Organisation, Bengaluru, India  | Jan 2023-Jun 2023 |
| <ul style="list-style-type: none"><li>• Engineered a <b>Temporal Graph Neural Network (GNN)</b>, leveraging continuous temporal data and node features to predict future user interactions on online platforms, increasing model accuracy by 2% over current benchmarks <a href="#">[Paper]</a>.</li><li>• Developed and integrated <b>Incremental BERT (iBERT)</b> with Temporal GNN to capture semantic drift and enhance real-time semantic understanding of evolving text data, reducing data processing time by 40%.</li><li>• Achieved 3.19 perplexity (6% better than SOTA) in masked language modeling, published in <b>Springer ICPR 2024</b> <a href="#">[Paper]</a>.</li></ul>  |                   |
| <b>Data Science Intern</b>   Merkle Inc., Bengaluru, India   | May 2022-Jul 2022 |
| <ul style="list-style-type: none"><li>• Led a team of 4 in developing <b>predictive models</b> (XGBoost, LightGBM, LSTM) for revenue optimization by transforming transactional data, applying <b>SQL indexing</b> on 10M+ records, and leveraging <b>LLM-based embeddings</b> to cluster product descriptions for segment-specific targeting, resulting in a 10% increase in campaign profitability.</li><li>• Processed <b>16M+</b> rows of Home Depot sales data using <b>PySpark</b>, improving query performance by 40% through advanced data handling techniques, generating actionable pricing insights.</li></ul>  |                   |

## PROJECTS & PUBLICATIONS

- **CoveredAI – Health Insurance Analysis App** [\[Code\]](#) : Built a **full-stack** AI-powered app using **React**, **Flask (RESTful APIs)**, **LangChain**, and **OpenAI GPT** to analyze, summarize, and compare health insurance documents. Integrated **RAG** (semantic search + chunking via FAISS) for natural language Q&A and plan comparisons. Enabled PDF/DOCX uploads, secure **Google OAuth**, and exportable reports, with plans to integrate privacy-preserving features like PII/PHI redaction.
- **Multimodal Conversation Derailment Detection** [\[Paper\]](#): Built a hierarchical **transformer** combining **BERT**, **Faster R-CNN**, and **GRU** for multimodal Reddit thread modeling, integrating text and visual cues. Achieved 71% accuracy and 78% AUC, outperforming text-only baselines by 6% in conversational derailment detection.
- **Legal Query AI Assistant** [\[Code\]](#) : Built an AI assistant using LLMs (**OpenAI GPT/LLaMA**) and **RAG** to deliver accurate legal query responses. Combined vector-based retrieval with semantic understanding and deployed a lightweight **Flask** interface for real-time contextual Q&A.
- **COVID-19 X-ray Detection** [\[Code\]](#) : Built a **CNN model** on **3-class X-ray dataset** (Normal, Pneumonia, COVID-19), achieving 95.3% training and 89.5% validation accuracy. Deployed a **Flask app** for real-time COVID detection from uploaded X-rays.
- **Privacy-Preserving LLM Evaluation** [\[Paper\]](#) : Analyzed **GPT-3.5**, **GPT-4**, and Turbo models for **PHI/PII leakage** on synthetic healthcare and hiring datasets. Achieved 60–99% privacy reduction while retaining >85 BLEU score, supporting **HIPAA/GDPR-compliant** LLM use.